Howard Hung-Mo Hsieh

Los Angeles, CA | hungmohsieh@gmail.com | linkedin.com/in/hung-mo-hsieh | github.com/howardhsieh

Education

University of Southern California — Los Angeles, CA
M.S. in Cyber Security Engineering, GPA 3.67/4.00 — Aug 2025 – Jun 2027 (expected)
- Coursework: Security Systems, Computer Systems Assurance, Security & Privacy, Computer Forensics

National Chung-Hsing University — Taichung, Taiwan
B.S. in Computer Science, GPA 3.70/4.30 — Sep 2019 – Jun 2023

Experience

Branchmore — May 2026 – Aug 2026
Security Engineering Intern — Remote, USA
- Strengthened cloud security by moving AWS SSH to Tailscale zero trust, cutting public ingress to zero
- Enforced least-privilege, just-in-time AWS access controls and wrote the team training guide
- Collaborated with engineers to gate high-blast-radius Terraform behind manual approval in GitHub Actions
- Deployed GuardDuty, Inspector, Prowler, and Semgrep for vulnerability scanning and monitoring
- Pen-tested Go and Elixir services and reported path traversal, command injection, and leaked secrets
- Closed 75 tickets over 4+ AWS accounts, documenting each with a verification runbook
- Coordinated a blue/green VPC migration of dev cloud infrastructure, announcing downtime early

GMVR Lab, National Chung-Hsing University — Apr 2024 – Mar 2025
Research Assistant — Taichung, Taiwan
- Co-authored an IEEE Access (2025) paper on reversible data hiding with lossless image restoration
- Implemented the algorithm in Python and OpenCV, raising embedding capacity by 15%

Systex Corporation — Aug 2023 – Apr 2024
Security Engineer — Taipei, Taiwan
- Managed Splunk SIEM for 15 enterprise clients, onboarding 5+ TB/day of Linux logs
- Developed 70+ SPL correlation analysis rules and dashboards, reducing false positive alerts by 30%
- Analyzed 20+ alerts/week and led incident response for clients

Projects

agent-policy-gateway — GitHub
- Built a security gateway enforcing access controls on LLM tool calls (MCP, OpenAI, Anthropic)
- Designed information-flow control in Python that blocks indirect prompt-injection exfiltration
- Tested with AgentDojo benchmarking, cutting attack success from 100% to 0% in 588 attacks
- Achieved 100% utility at 0% compromise on 90 workflows, outperforming Progent and FIDES
- Automated 1,400+ tests in GitHub Actions CI workflows and shipped v0.1 to PyPI

TraceSig — GitHub
- Engineered a Sigma-style rule format and parsing engine that automates attack detection in LLM traces
- Wrote 24 YAML rules across three packs for Claude Code and gateway logs; shipped v0.2.1 to PyPI

preflight | **blastradius** — GitHub | GitHub
- Built two Go CLIs: a scanner for auto-run code in untrusted repos and an offline credential inventory

Skills

Security: zero trust, detection engineering (Splunk SPL, Sigma), incident response, threat modeling, SAST/SCA (gosec, Sobelow, gitleaks), CSPM, AI/LLM security

Cloud & DevOps: AWS (IAM, SCPs, VPC, EC2, Lambda, CloudTrail, GuardDuty), Terraform, GitHub Actions, Linux

Languages: Python, Go, Bash, SQL, C/C++, Java

Certifications: Splunk Core Certified Power User; CompTIA SecAI+ (in progress, Oct 2026)